

Lecture 6

Sorting lower bounds and $O(n)$ -time sorting

Announcements


- Remind me about a break around 11:20!
- Remind me to repeat the question!
- New additional resource: last year's lecture notes
 - **NOT required reading**, totally optional
 - Intermediate level of rigor (between the slides and CLRS)
 - Based on Virginia Williams's lecture notes from an earlier offering of CS161
 - NOT being maintained for this quarter
 - (In particular, not valid for bug bounty points 😊)
 - Posted on the website in case you find them helpful.

Sorting

- We've seen a few $O(n \log(n))$ -time algorithms.
 - MERGESORT has worst-case running time $O(n \log(n))$
 - QUICKSORT has expected running time $O(n \log(n))$

Can we do better?


Depends on who
you ask...





An $O(1)$ -time algorithm for sorting: StickSort

- Problem: sort these n sticks by length.



- Now they are sorted this way.

- Algorithm:
 - Drop them on a table.

That may have been unsatisfying

- But StickSort does raise some important questions:

- What is our model of computation?

- Input: array
- Output: sorted array
- Operations allowed: comparisons

-vs-

- Input: sticks
- Output: sorted sticks in vertical order
- Operations allowed: dropping on tables

- What are reasonable models of computation?

Today: two (more) models




- Comparison-based sorting model
 - This includes MergeSort, QuickSort, InsertionSort
 - We'll see that any algorithm in this model must take at least $\Omega(n \log(n))$ steps.



- Another model (more reasonable than the stick model...)
 - BucketSort and RadixSort
 - Both run in time $O(n)$

Comparison-based sorting



Comparison-based sorting algorithms




😊 is shorthand for
“the first thing in the input list”

Want to sort these items.

There's some ordering on them, but we don't know what it is.

Is  bigger than  ?



Algorithm

YES

The algorithm's job is to
output a correctly sorted
list of all the objects.



There is a **genie** who knows what
the right order is.

The genie can answer YES/NO
questions of the form:
is [this] bigger than [that]?

All the sorting algorithms we have seen work like this.

eg, QuickSort:



Is 7 bigger than 5 ?

YES

Is 6 bigger than 5 ?


YES


Is 3 bigger than 5 ?

NO



etc.





Lower bound of $\Omega(n \log(n))$.

- Theorem:

- Any **deterministic comparison-based sorting algorithm** must take $\Omega(n \log(n))$ steps.
- Any **randomized comparison-based sorting algorithm** must take $\Omega(n \log(n))$ steps in expectation.

*This covers all the
sorting algorithms
we know!!!*

- How might we prove this?

1. Consider all comparison-based algorithms, one-by-one, and analyze them.

2. Don't do that.

Instead, argue that all comparison-based sorting algorithms give rise to a **decision tree**.
Then analyze decision trees.

Decision trees



Sort these three things.



YES

NO

etc...



YES

NO




YES

NO



Decision trees

- Internal nodes correspond to yes/no questions.
- Each internal node has two children, one for “yes” and one for “no.”
- Leaf nodes correspond to outputs.
 - In this case, all possible orderings of the items.
- Running an algorithm on a particular input corresponds to a particular path through the tree.



Comparison-based algorithms look like decision trees.

Pivot!



Smiley face \leq Fire truck ?

YES

NO




etc...

Coffee cup \leq Fire truck ?

YES

NO




Pivot!

Now
recurse
on R

Smiley face \leq Coffee cup ?

YES

NO



Return



Return



Return



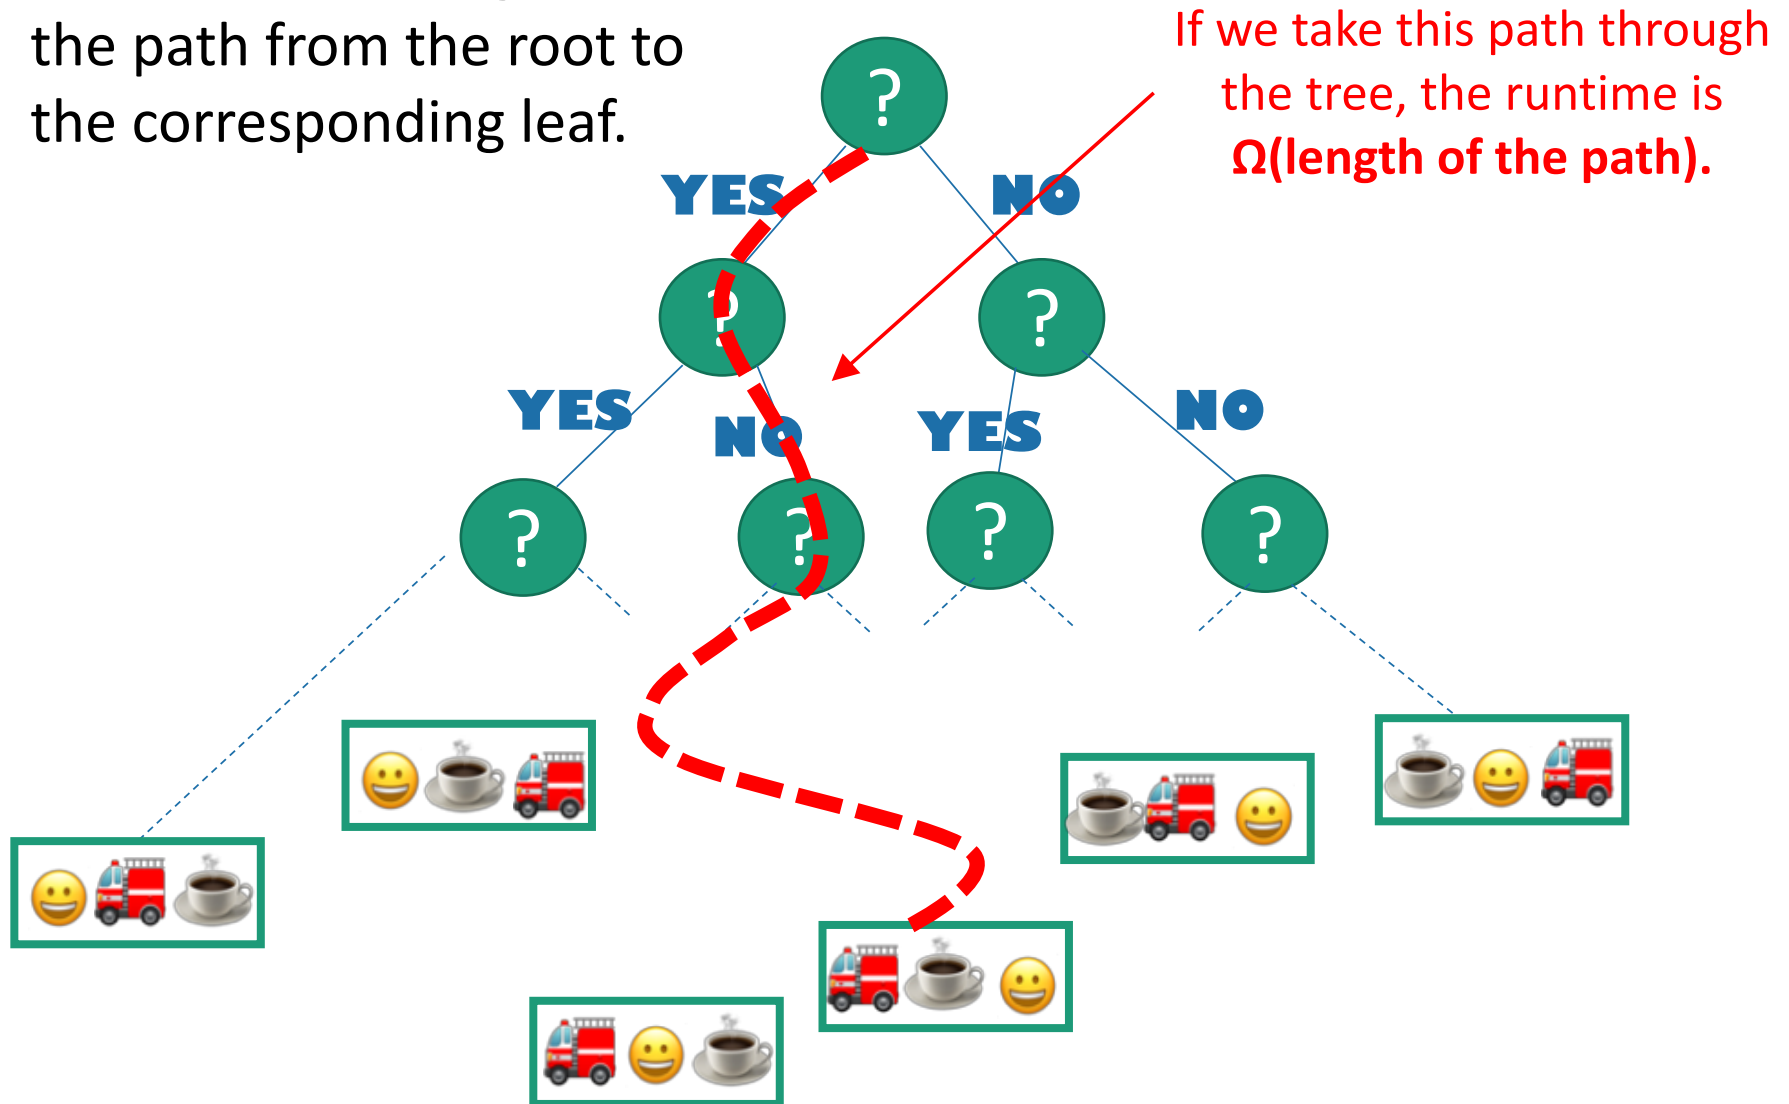
Example: Sort these
three things using
QuickSort.

Then we're done
(after some base-
case stuff)

In either case, we're done
(after some base case stuff and
returning recursive calls).

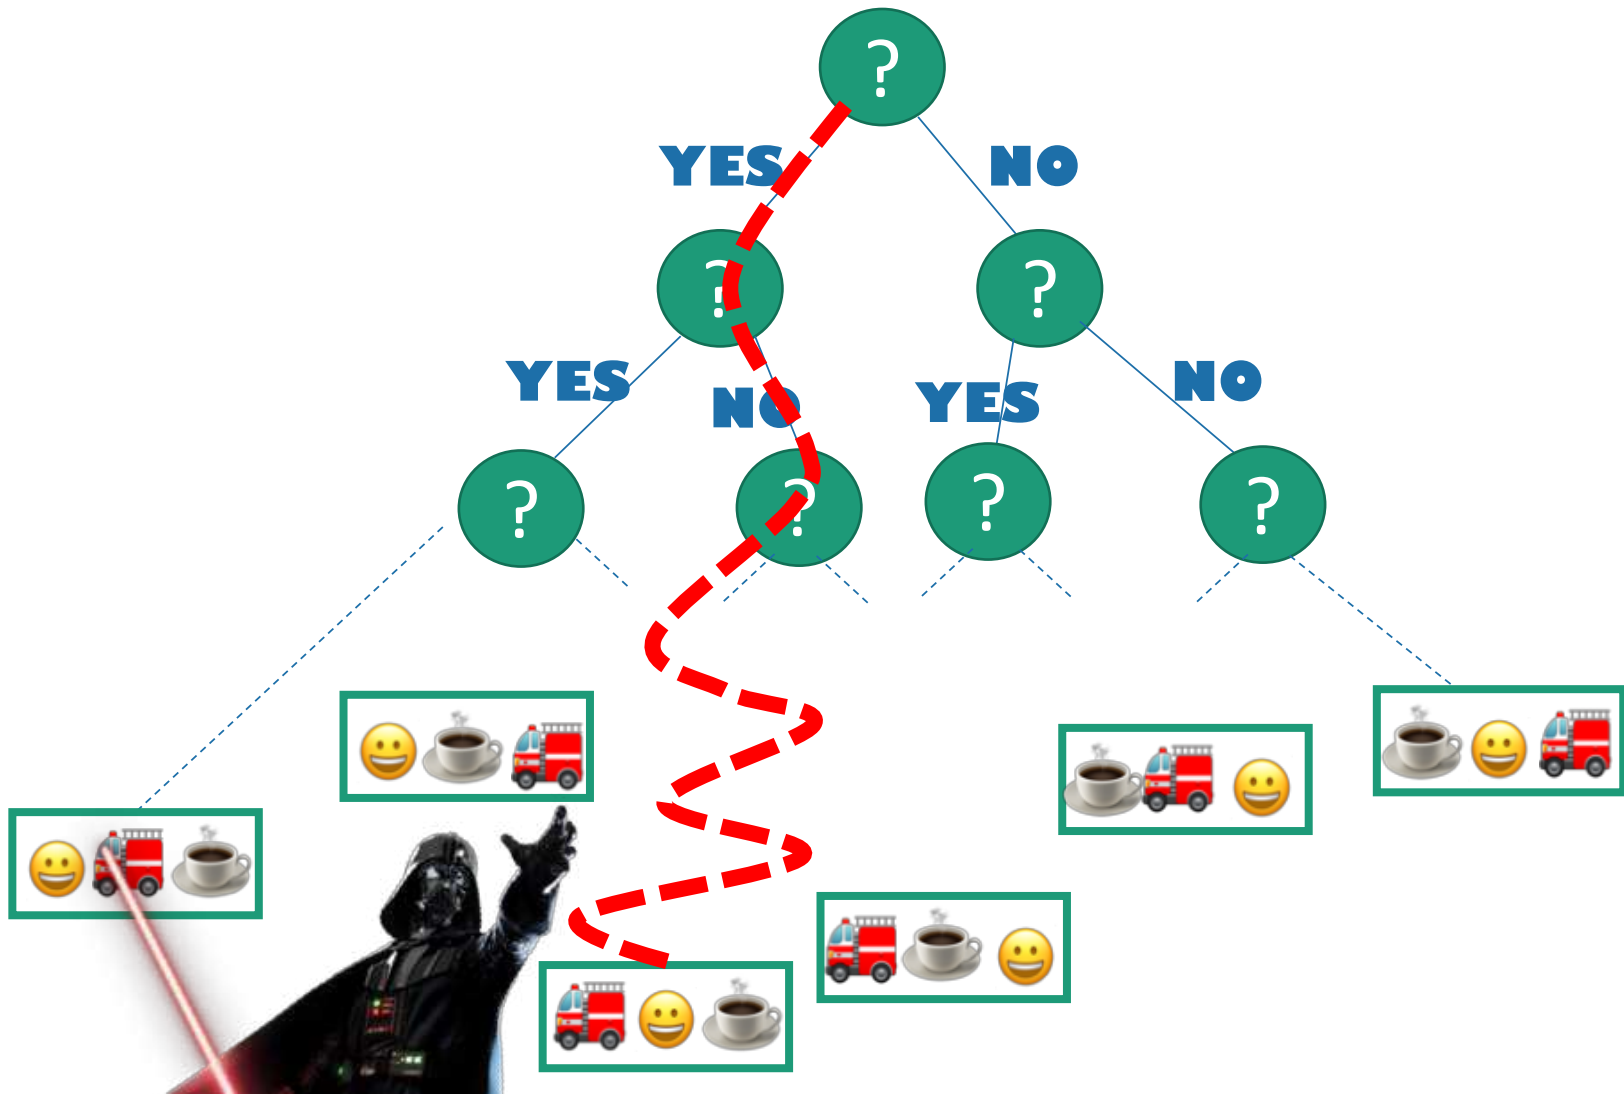
Q: What's the runtime on a particular input?

A: At least the length of the path from the root to the corresponding leaf.



Q: What's the worst-case runtime?

A: At least $\Omega(\text{length of the longest path})$.




How long is the longest path?



being sloppy about
floors and ceilings!

We want a statement: in all such trees,
the longest path is at least _____




- This is a binary tree with at least $n!$ leaves.
- The shallowest tree with $n!$ leaves is the completely balanced one, which has depth $\log(n!)$.
- So in all such trees, the longest path is at least $\log(n!)$.

- $n!$ is about $(n/e)^n$ (Stirling's approx.*).
- $\log(n!)$ is about $n \log(n/e) = \Omega(n \log(n))$.

Conclusion: the longest path
has length at least $\Omega(n \log(n))$.

*Stirling's approximation is a bit more complicated than this, but this is good enough for the asymptotic result we want.

Lower bound of $\Omega(n \log(n))$.



- **Theorem:**

- Any deterministic comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps.

- **Proof recap:**

- Any deterministic comparison-based algorithm can be represented as a decision tree with $n!$ leaves.
- The worst-case running time is at least the depth of the decision tree.
- All decision trees with $n!$ leaves have depth $\Omega(n \log(n))$.
- So any comparison-based sorting algorithm must have worst-case running time at least $\Omega(n \log(n))$.

Aside:

What about randomized algorithms?

- For example, QuickSort?

- Theorem:

- Any randomized comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps in expectation.



- Proof:

- see reading posted on website
 - (Avrim Blum's notes)
- (same ideas as deterministic case)


Try to prove this
yourself!



\end{Aside}

Ollie the over-achieving ostrich

So that's bad news



- Theorem:

- Any deterministic comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps.

- Theorem:

- Any randomized comparison-based sorting algorithm must take $\Omega(n \log(n))$ steps in expectation.

On the bright side,
MergeSort is optimal!

- This is one of the cool things about lower bounds like this:
we know when we can declare victory!



But what about StickSort?

- StickSort can't be implemented as a comparison-based sorting algorithm. So these lower bounds don't apply.
- But StickSort was kind of silly.

Can we do better?

- Is there be another model of computation that's **less silly** than the StickSort model, in which we can **sort faster** than $n\log(n)$?

Especially if I have
to spend time
cutting all those
sticks to be the
right size!



Beyond comparison-based sorting algorithms



Another model of computation

- The items you are sorting have **meaningful values**.



instead of



Pre-lecture exercise

- How long does it take to sort n people by their month of birth?
- [discussion]



1 (Jan)



1 (Jan)




4 (Apr)



5 (May)

Another model of computation

- The items you are sorting have **meaningful values**.



instead of




Why might this help?



Implement the buckets as linked lists. They are first-in, first-out. This will be useful later.

BucketSort:

Note: this is a simplification of what CLRS calls “BucketSort”



Concatenate
the buckets!

SORTED!


In time $O(n)$.

Assumptions

- Need to be able to know what bucket to put something in.
 - We assume we can evaluate the items directly, not just by comparison
- Need to know what values might show up ahead of time.

2	12345	13	2^{1000}	50	100000000	1
---	-------	----	------------	----	-----------	---



- Need to assume there are not too many such values.




RadixSort

- For sorting integers up to size M
 - or more generally for lexicographically sorting strings
- Can use less space than BucketSort
- Idea: BucketSort on the least-significant digit first, then the next least-significant, and so on.


Step 1: BucketSort on least significant digit



Step 2: BucketSort on the 2nd least sig. digit



Step 3: BucketSort on the 3rd least sig. digit



It worked!!

Why does this work?

Original array:

21	345	13	101	50	234	1
----	-----	----	-----	----	-----	---

Next array is sorted by the first digit.

50	21	101	1	13	234	345
----	----	-----	---	----	-----	-----

Next array is sorted by the first two digits.

101	01	13	21	234	345	50
-----	----	----	----	-----	-----	----

Next array is sorted by all three digits.

001	013	021	050	101	234	345
-----	-----	-----	-----	-----	-----	-----

Sorted array

To prove this is correct...

- What is the inductive hypothesis?

Original array:

21	345	13	101	50	234	1
----	-----	----	-----	----	-----	---

Next array is sorted by the first digit.

50	21	101	1	13	234	345
----	----	-----	---	----	-----	-----


Next array is sorted by the first two digits.

101	01	13	21	234	345	50
-----	----	----	----	-----	-----	----

Next array is sorted by all three digits.

001	013	021	050	101	234	345
-----	-----	-----	-----	-----	-----	-----

Sorted array



Think-Pair-Share Terrapins

RadixSort is correct

- Inductive hypothesis:
 - After the k 'th iteration, the array is sorted by the first k least-significant digits.
- Base case:
 - “Sorted by 0 least-significant digits” means not sorted, so the IH holds for $k=0$.
- Inductive step:
 - TO DO
- Conclusion:
 - The inductive hypothesis holds for all k , so after the last iteration, the array is sorted by all the digits. Hence, it's sorted!

Inductive step


Inductive hypothesis:

After the k 'th iteration, the array is sorted by the first k least-significant digits.

- Need to show: if IH holds for $k=i-1$, then it holds for $k=i$.
 - Suppose that after the $i-1$ 'st iteration, the array is sorted by the first $i-1$ least-significant digits.
 - Need to show that after the i 'th iteration, the array is sorted by the first i least-significant digits.

EXAMPLE: $i=2$

IH: this array is sorted by first digit.



Want to show: this array is sorted by 1st and 2nd digits.

Proof sketch...

proof on next (skipped) slide

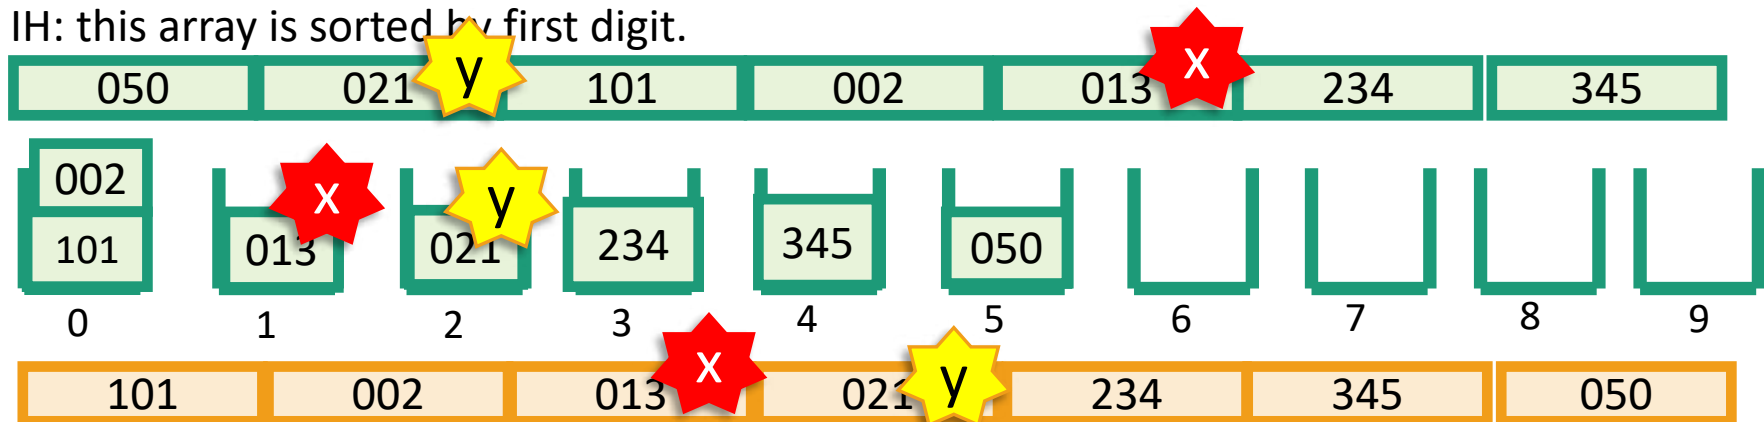
Want to show: after the i 'th iteration, the array is sorted by the first i least-significant digits.

- Write $x=[x_dx_{d-1}...x_2x_1]$ and $y=[y_dy_{d-1}...y_2y_1]$
- Suppose $[x_ix_{i-1}...x_2x_1] < [y_iy_{i-1}...y_2y_1]$.
- Want to show that x appears before y at end of i 'th iteration.
- CASE 1: $x_i < y_i$
 - x is in an earlier bucket than y .



EXAMPLE: $i=2$

IH: this array is sorted by first digit.



Want to show: this array is sorted by 1st and 2nd digits.

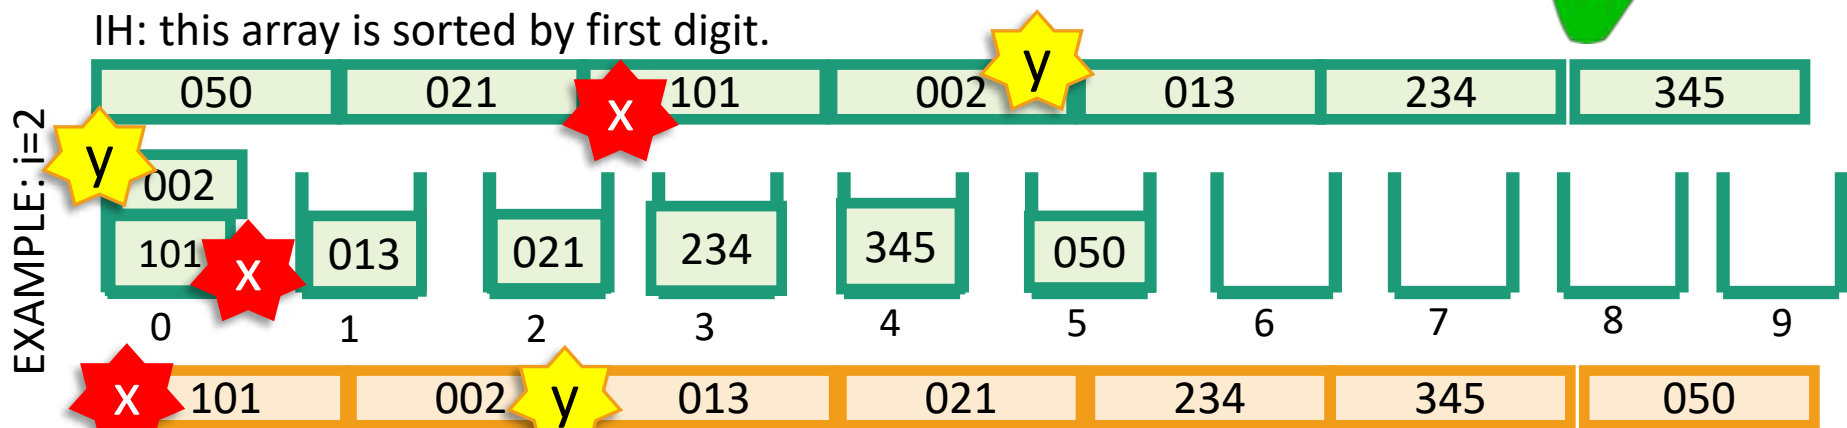
Proof sketch...

proof on next (skipped) slide

Want to show: after the i 'th iteration, the array is sorted by the first i least-significant digits.

- Let $x=[x_dx_{d-1}...x_2x_1]$ and $y=[y_dy_{d-1}...y_2y_1]$ be any x,y so that $[x_ix_{i-1}...x_2x_1] < [y_iy_{i-1}...y_2y_1]$.
- Want to show that x appears before y at end of i 'th iteration.
- CASE 1: $x_i < y_i$**
 - x is in an earlier bucket than y .
- CASE 2: $x_i = y_i$**
 - x and y in same bucket, but x was put in the bucket first.

Aka, we want to show that for any x and y so that x belongs before y , we put x before y .



Want to show: this array is sorted by 1st and 2nd digits.

Want to show: after the i 'th iteration, the array is sorted by the first i least-significant digits.

SLIDE SKIPPED
IN CLASS. Here
for reference.

- Write $x = [x_d x_{d-1} \dots x_2 x_1]$ and $y = [y_d y_{d-1} \dots y_2 y_1]$
- Suppose $[x_i x_{i-1} \dots x_2 x_1] < [y_i y_{i-1} \dots y_2 y_1]$.
- Want to show that x appears before y at end of i 'th iteration.
- CASE 1: $x_i < y_i$.
 - x appears in an earlier bucket than y , so x appears before y after the i 'th iteration.
- CASE 2: $x_i = y_i$.
 - x and y end up in the same bucket.
 - In this case, $[x_{i-1} \dots x_2 x_1] < [y_{i-1} \dots y_2 y_1]$, so by the inductive hypothesis, x appeared before y after $i-1$ 'st iteration.
 - Then x was placed into the bucket before y was, so it also comes out of the bucket before y does.
 - Recall that the buckets are FIFO queues.
 - So x appears before y in the i 'th iteration.

Inductive step


Inductive hypothesis:

After the k 'th iteration, the array is sorted by the first k least-significant digits.

- Need to show: if IH holds for $k=i-1$, then it holds for $k=i$.
 - Suppose that after the $i-1$ 'st iteration, the array is sorted by the first $i-1$ least-significant digits.
 - Need to show that after the i 'th iteration, the array is sorted by the first i least-significant digits.




IH: this array is sorted by first digit.



Want to show: this array is sorted by 1st and 2nd digits.

RadixSort is correct

- Inductive hypothesis:
 - After the k 'th iteration, the array is sorted by the first k least-significant digits.
- Base case:
 - “Sorted by 0 least-significant digits” means not sorted, so the IH holds for $k=0$.
- Inductive step:
 - TO DO 
- Conclusion:
 - The inductive hypothesis holds for all k , so after the last iteration, the array is sorted by all the digits. Hence, it's sorted!


What is the running time? for RadixSorting numbers base-10.

- Suppose we are sorting n d -digit numbers (in base 10).

e.g., $n=7$, $d=3$:

021	345	013	101	050	234	001
-----	-----	-----	-----	-----	-----	-----

1. How many iterations are there?
2. How long does each iteration take?
3. What is the total running time?



Think-Pair-Share Terrapins

What is the running time? for RadixSorting numbers base-10.

- Suppose we are sorting n d -digit numbers (in base 10).

e.g., $n=7$, $d=3$:

021	345	013	101	050	234	001
-----	-----	-----	-----	-----	-----	-----

1. How many iterations are there?


- d iterations

2. How long does each iteration take?

- Time to initialize 10 buckets, plus time to put n numbers in 10 buckets. $O(n)$.

3. What is the total running time?


- $O(nd)$



Think-Pair-Share Terrapins

This doesn't seem so great

- To sort n integers, each of which is in $\{1, 2, \dots, n\}$...
- $d = \lfloor \log_{10}(n) \rfloor + 1$
 - For example:
 - $n = 1234$
 - $\lfloor \log_{10}(1234) \rfloor + 1 = 4$
 - More explanation on next (skipped) slide.
- Time = $O(nd) = O(n \log(n))$.
 - Same as MergeSort!



Aside: why $d = \lfloor \log_{10}(n) \rfloor + 1$?

Slide
skipped
in class

- When we write a number $x = [x_d x_{d-1} \dots x_1]$ base 10, that means:
$$x = x_1 + x_2 \cdot 10 + \dots + x_{d-1} \cdot 10^{d-2} + x_d \cdot 10^{d-1}$$

where $x_i \in \{0, 1, \dots, 9\}$

- Suppose that $x_d \neq 0$. Then we have

- $x \geq x_d \cdot 10^{d-1}$

Since x is bigger than just the last term in that sum!

- $\log_{10}(x) + 1 - \log_{10}(x_d) \geq d$

(take logs of both sides and rearrange)

- $\log_{10}(x) + 1 > d$

$\log_{10}(x_d) > 0$ since $x_d > 0$

- $\lfloor \log_{10}(n) \rfloor + 1 \geq d$

Since d is an integer

- On the other hand, we also have

- $x < (x_d + 1) \cdot 10^{d-1}$

Since if $x \geq (x_d + 1) \cdot 10^{d-1}$ then the d 'th digit would have been $x_d + 1$ instead of x_d

- $\log_{10}(x) + 1 - \log_{10}(x_d + 1) < d$

(take logs of both sides and rearrange)

- $\log_{10}(x) < d$

$\log_{10}(x_d + 1) \leq 1$ since $x_d < 10$

- $\lfloor \log_{10}(n) \rfloor + 1 \leq d$

Since d is an integer

Can we do better?

- RadixSort base 10 doesn't seem to be such a good idea...
- But what if we change the base? (Let's say base r)
- We will see there's a trade-off:
 - Bigger r means more buckets
 - Bigger r means fewer digits




Example: base 100

Original array:



21	345	13	101	50	234	1
----	-----	----	-----	----	-----	---

Example: base 100


Original array:




100 buckets:



Example: base 100



100 buckets:



Sorted!

Example: base 100

Original array

0021	0345	0013	0101	0050	0234	0001
------	------	------	------	------	------	------

0101	0001	0013	0021	0234	0345	0050
------	------	------	------	------	------	------

0001	0013	0021	0050	0101	0234	0345
------	------	------	------	------	------	------

Sorted array

Base 100:

- $d=2$, so only 2 iterations.
- 100 buckets

vs.

Base 10:

- $d=3$, so 3 iterations.
- 10 buckets

Bigger base means more buckets but fewer iterations.

General running time of RadixSort



- Say we want to sort:
 - n integers,
 - maximum size M ,
 - in base r .
- Number of iterations of RadixSort:
 - Same as number of digits, base r , of an integer x of max size M .
 - That is $d = \lfloor \log_r(M) \rfloor + 1$
- Time per iteration:
 - Initialize r buckets, put n items into them
 - $O(n + r)$ total time.
- Total time:
 - $O(d \cdot (n + r)) = O((\lfloor \log_r(M) \rfloor + 1) \cdot (n + r))$

Convince yourself that this is the right formula for d .



Trade-offs

- Given n , M , how should we choose r ?
- Looks like there's some sweet spot:



A reasonable choice: $r=n$

- Running time:

$$O((\lfloor \log_r(M) \rfloor + 1) \cdot (n + r))$$

Intuition: balance n and r here.

- Choose $n=r$:

$$O(n \cdot (\lfloor \log_n(M) \rfloor + 1))$$

Choosing $r = n$ is pretty good. What choice of r optimizes the asymptotic running time? What if I also care about space?



Ollie the over-achieving ostrich

Running time of RadixSort with $r=n$

- To sort n integers of size at most M , time is

$$O\left(n \cdot (\lfloor \log_n(M) \rfloor + 1)\right)$$

- So the running time (in terms of n) depends on how big M is in terms of n :
 - If $M \leq n^c$ for some constant c , then this is $O(n)$.
 - If $M = 2^n$, then this is $O\left(\frac{n^2}{\log(n)}\right)$
- The number of buckets needed is $r=n$.


What have we learned?

You can put any
constant here
instead of 100.

- RadixSort can sort n integers of size at most n^{100} in time $O(n)$, and needs enough space to store $O(n)$ integers.
- If your integers have size much much bigger than n (like 2^n), maybe you shouldn't use RadixSort.
- It matters how we pick the base.



Recap

- How difficult sorting is depends on the model of computation.
- How reasonable a model of computation is is up for debate.
- Comparison-based sorting model
 - This includes MergeSort, QuickSort, InsertionSort
 - Any algorithm in this model must use at least $\Omega(n \log(n))$ operations. 😞
 - But it can handle arbitrary comparable objects. 😊
- If we are sorting small integers (or other reasonable data):
 - BucketSort and RadixSort 
 - Both run in time $O(n)$ 😊
 - Might take more space and/or be slower if integers get too big 😞



Next time

- Binary search trees!
- Balanced binary search trees!

Before next time

- Pre-lecture exercise for Lecture 7
 - Remember binary search trees?

**CHUCK NORRIS
QUICKSORTS STICKS**



IN TIME O(1)